
NEWS BROADCAST ANALYSIS

Christine K. C. Cheng

1 Shot detection

Sum of absolute differences

A score is assigned to each frame by calculating the sum of the absolute differences between consecutive frames for every pixel. This value is then normalized by the size of the frame. A threshold is selected through empirical experiments before score calculations. When the score of a frame is greater than the threshold, a shot change is declared. This method works quite well with simple videos but it is not robust against movements and changes in lighting.

Histogram differences

Each frame is converted into a gray-scale image. A histogram with 256 bins, representing all the possible values of a pixel, is created for each frame. Then, a score is assigned by calculating the sum of the absolute differences between histograms of consecutive frames.

Figure 1: Sum of absolute differences scores of clip 1




Figure 2: Histogram differences scores of clip 1




Figure 3: Sum of absolute differences scores of clip 2




Figure 4: Histogram differences scores of clip 2




Figure 5: Sum of absolute differences scores of clip 3





Figure 6: Histogram differences scores of clip 3



The relevant code is in `shot.py`. To get the graphs of shot detection, run the command below.

```
python3 run.py shot_detection -t <type> -i <path to frames>
```

2 Logo detection

Template matching is an object detection algorithm which is translation invariant but not scale or rotation invariant. As we are detecting a logo, we can assume that the target of detection will be in a known orientation. Due to the possibility that there may be multiple occurrences of the logo in a frame, we cannot simply match SIFT features between the logo template and a frame. Template matching is run on templates of different sizes because the size of the logo in a frame is unknown. Then, a score is calculated for all the matches by normalized cross-correlation. The normalized version is chosen because brighter patches will not have a higher score. Also, the score obtained will be in the range $[0, 1]$ and this makes choosing a threshold more intuitive. As the logo of the template and the one in the frame might be of different, possibly due to different resolutions or styles, a looser threshold is first used to filter out the irrelevant matches.

For each match, if its score is greater than the loose threshold, it is kept. Then, the algorithm checks if that particular match is a slight translation of a match we have already decided to keep. If it is, the match is discarded. This prevents having multiple boxes around one logo. Afterwards, SIFT descriptors are calculated for the remaining matches and a score is calculated using feature matching and Lowe's ratio test with the logo template. If a match's score is above a tighter threshold, it is declared as a match of the template and a box is put around the match.

The logo detection was the most difficult part of the project for me. I originally only did one pass with either normalized cross-correlation or SIFT feature matching. However, this led to fairly poor results, with the algorithm often unable to detect multiple logos and including many irrelevant matches. Therefore, the two passes approach, first with a looser threshold using normalized cross-correlation and then a tighter threshold with SIFT feature matching, is used. As can be seen in Figure 8, the algorithm can detect multiple logos.

Figure 7: Logo detection on frame 104 of clip 1



Figure 8: Logo detection on frame 52 of clip 1



The relevant code is in `logo.py`. To run logo detection, run the command below.

```
python3 run.py logo_detection -i <input directory> -o <output directory>
-d <logo path> -t <min threshold>
```

3 Face detection and tracking

There are 260 images in the female and male classes respectively. Each image is accompanied by a .mat file specifying the coordinates of the left eye, right eye, nose and mouth. The following rules are used to crop the images in order to obtain the faces.

$$\begin{aligned} start_x &= left\ eye_x - 0.5 \times (right\ eye_x - left\ eye_x) \\ end_x &= right\ eye_x + 0.5 \times (right\ eye_x - left\ eye_x) \\ start_y &= eyes_y - (mouth_y - eyes_y) \\ end_y &= mouth_y + (mouth_y - eyes_y) \end{aligned} \quad (1)$$

The relevant code for face cropping is in `crop_images()` in `face.py`.

The initial attempt to detect faces is to use skin detection - trying to filter out skin in images. Figure 9 and Figure 10 show the color distributions of faces in RGB and HSV color spaces. The HSV color space has narrower distributions, especially with hue. Although this method works sometimes as seen in Figure 11, it is not successful in general. It fails to detect a large area of the face of the man on the right in Figure 12 and includes a lot of the background. This model is especially poor when other things in the frame are very similar to human skin tone.

Figure 9: RGB distribution of female training images

Figure 10: HSV distribution of female training images


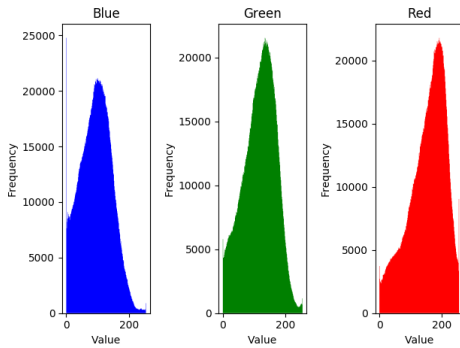



Figure 11: HSV colour detection on frame 160 of clip 1

Figure 12: HSV colour detection on frame 50 of clip 1



The relevant code for HSV face detection is in `visualize_distributions()` and `face_detection_hsv()` in `face.py`.

I ended up using `cv2.CascadeClassifier` for detecting faces. The full name of this classifier is Haar feature-based cascade classifier for object detection. It is a OpenCV pre-trained classifier for face stored in an XML file. It works fairly well. It has no trouble detecting multiple people in a frame or people of colour, like in Figure 13. However, it is sometimes unable to detect faces in a certain position. The detector was unable to detect the person on the left in any of the frames he appeared in that particular position in Figure 14.

After obtaining the faces in a frame, for each face in the current frame, the SIFT descriptors are found. Then, they are matched to the SIFT descriptors of each faces in the previous frame using feature matching and Lowe's ratio test. Then, a score is obtained from the number of matches. If the score is above a set threshold, the algorithm declares that the face we are looking at is found in the previous frame and we will display the index assigned to that particular face in the previous frame. If the face is not found, we assign a new index to the face.

Figure 13: Face detection on frame 69 of clip 2



Figure 14: Face detection on frame 110 of clip 1



4 Gender classification

90% of the images (234 images from each class) are used for training, whereas the other 10% (26 images from each class) are used for testing the accuracy of the model.

SVM

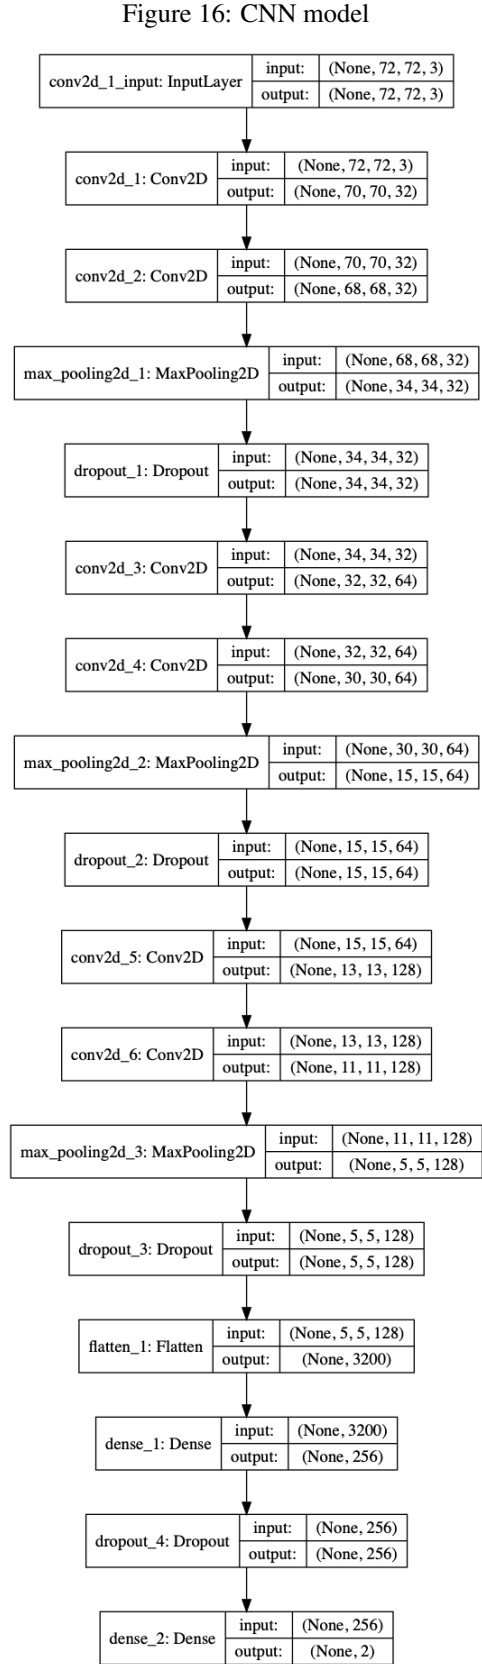
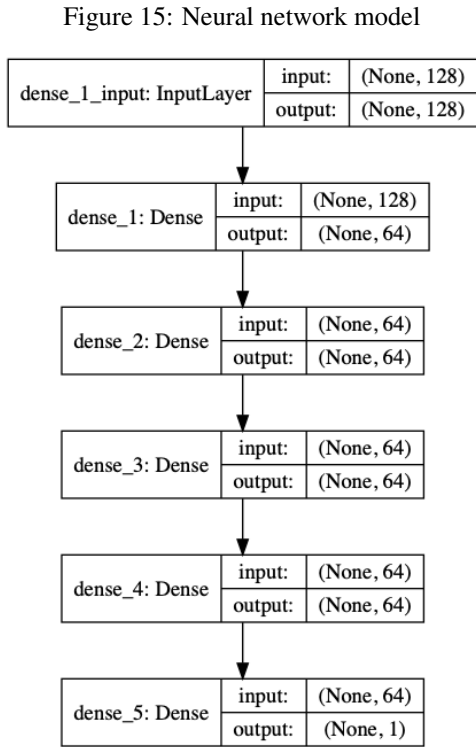
The SIFT descriptors of the training images are passed into the SVM model for training. For each detected face, the SIFT descriptors are extracted and fed to the trained SVM model. Then, a prediction for each descriptor is obtained. If more descriptors are predicted as female than male, the image is classified as female. If more descriptors are predicted as male than female, the image is classified as male. If there are equal number of descriptors being predicted as both female and male, the image is then classified as unknown.

Neural network

The same as SVM, except with a neural network model instead. The model uses a binary crossentropy loss, adam for optimization, and accuracy as the metric.

CNN

All the training and testing images are padded with black borders to obtain a square shape and then resized to be 72 pixels by 72 pixels. The training images are then passed to the CNN model shown in Figure 16 for training. The model uses a binary crossentropy loss, stochastic gradient descent for optimization, and accuracy as the metric. Each detected face is padded to obtain a square shape. Then, the image is resized to be 72 pixels by 72 pixels. The resized image is then passed to the trained cnn model and a category prediction is obtained.



The relevant code for training any of the three models is in `train_model()` in `face.py`. To train a gender classification model, run the command below.

```
python3 run.py train -m <model path after training>
-c <classification model (SVM, NN_SIFT, or CNN)>
```

The relevant code for face detection (including gender classification and face tracking) is in `face_detection()` in `face.py`. Run the command below.

```
python3 run.py face_detection -i <input directory> -o <output directory>
-c <classification model (SVM, NN_SIFT, or CNN)> -m <trained model path>
```

Performance

Table 1 shows the test accuracies of the three models. As expected, CNN performed poorly, achieving an accuracy that is equivalent to random guesses, due to the very small training data size of 468.

Table 1: Gender classification performance

Model	Description	Accuracy on test set
SVM	Using SIFT descriptors of faces	100.00%
Neural network	Using SIFT descriptors of faces	92.30%
CNN	Using cropped and resized faces	50.00%

5 References

Video shot boundary detection based on color histogram

Wikipedia - Shot transition detection

6 Code

Listing 1: shot.py

```
1 import cv2
2 import matplotlib.pyplot as plt
3 import numpy as np
4 import os
5
6 def shot_detection(input_dir, type):
7     scores = []
8     img_names = [img for img in os.listdir(input_dir) if img.endswith(".jpg")]
9     # Sort images by name in ascending order
10    img_names.sort(key=lambda img: int("".join(filter(str.isdigit, img))))
11    start_idx = int(img_names[0][:img_names[0].find(".jpg")])
12    end_idx = int(img_names[-1][:img_names[-1].find(".jpg")])
13
14    # Sum of absolute differences
15    if type == "SAD2":
16        prev_img = None
17        for i in range(len(img_names)):
18            # Default type is numpy.uint64
19            curr_img = cv2.imread(os.path.join(input_dir, img_names[i])).astype(np.int64)
20            if i == 0:
21                r, c, d = curr_img.shape
22                next_img = cv2.imread(os.path.join(input_dir, img_names[i + 1])).astype(np.
23                    ↪ int64)
24                score = np.sum(np.abs(curr_img - next_img))
25            elif i == len(img_names) - 1:
26                score = np.sum(np.abs(curr_img - prev_img))
27            else:
28                next_img = cv2.imread(os.path.join(input_dir, img_names[i + 1])).astype(np.
29                    ↪ int64)
30                score = 0.5 * np.sum(np.abs(curr_img - prev_img)) + 0.5 * np.sum(np.abs(
31                    ↪ curr_img - next_img))
32            scores.append(score)
33            prev_img = curr_img
34        x = np.arange(start_idx, end_idx + 1)
35        scores = np.array(scores) / (r * c * d)
36        title = "Sum of absolute differences"
37        new_filename = "output/" + input_dir.name + "_score_sad2.png"
38
39    elif type == "SAD":
40        prev_img = None
41        for i in range(len(img_names)):
42            # Default type is numpy.uint64
43            curr_img = cv2.imread(os.path.join(input_dir, img_names[i])).astype(np.int64)
44            if i == 0:
45                r, c, d = curr_img.shape
46                prev_img = curr_img
47                continue
48            score = np.sum(np.abs(curr_img - prev_img))
49            scores.append(score)
50            prev_img = curr_img
51        x = np.arange(start_idx + 1, end_idx + 1)
52        scores = np.array(scores) / (r * c * d)
53        title = "Sum of absolute differences"
54        new_filename = "output/" + input_dir.name + "_score_sad2.png"
55
56    # Histogram differences
```

```

54     elif type == "HD":
55         prev_histogram = None
56         for i in range(len(img_names)):
57             # Default type is numpy.uint8
58             curr_img = cv2.imread(os.path.join(input_dir, img_names[i]))#.astype(np.int64)
59             # Default type is numpy.uint8
60             curr_img_g = cv2.cvtColor(curr_img, cv2.COLOR_BGR2GRAY).astype(np.int16)
61             histogram = np.histogram(np.ravel(curr_img_g), bins=np.arange(-1, 256))
62             if i == 0:
63                 r, c = curr_img_g.shape
64                 prev_histogram = histogram[0]
65                 continue
66             score = np.sum(np.abs(histogram[0] - prev_histogram))
67             scores.append(score)
68             prev_histogram = histogram[0]
69
70             x = np.arange(start_idx + 1, end_idx + 1)
71             scores = np.array(scores) / (r * c)
72             title = "Histogram differences"
73             new_filename = "output/" + input_dir.name + "_score_hd.png"
74
75     else:
76         raise ValueError("Illegal type value")
77
78     print(len(x), len(scores))
79     f = plt.figure(figsize=(10, 5))
80     ax = f.gca()
81     ax.set_xticks(np.arange(start_idx, end_idx, 10))
82     plt.title(title)
83     plt.xlabel("Frame")
84     plt.ylabel("Score")
85     plt.plot(x, scores)
86     plt.grid()
87     f.savefig(new_filename)
88     print("Output saved to " + new_filename)

```

Listing 2: logo.py

```

1  import cv2
2  import imutils
3  import numpy as np
4  import os
5
6  import copy
7
8
9  # def logo_detection(input_dir, output_dir, logo_path):
10 # print("detect_logo")
11
12 # if not os.path.isdir(output_dir):
13 # os.mkdir(output_dir)
14
15 # template = cv2.imread(str(logo_path))
16 # template_g = cv2.cvtColor(template, cv2.COLOR_BGR2GRAY)
17
18 # w, h = template_g.shape[: -1]
19
20 # # Compute template of different sizes
21 # scales = np.linspace(0.1, 1.0, 25)[:: -1]
22 # templates = []

```



```

23 # ratios = []
24 # for scale in scales:
25 #     resized = imutils.resize(template_g, width=int(template_g.shape[1] * scale))
26 #     template_canny = cv2.Canny(resized, 50, 200)
27 #     templates.append(template_canny)
28 #     ratios.append(resized.shape[1] / float(template_g.shape[1]))
29
30 # for img_name in os.listdir(input_dir):
31 #     if not img_name.endswith(".jpg"):
32 #         continue
33 #     img = cv2.imread(os.path.join(input_dir, img_name))
34 #     img_g = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
35 #     img_canny = cv2.Canny(img_g, 50, 200)
36 #     cv2.imshow("img_canny", img_canny)
37 #     cv2.waitKey(0)
38
39 # found = None
40 # Loop through the templates from small to big
41 # for i in range(len(templates) - 1, -1, -1):
42 #     Stop when template is bigger than image
43 #     if img_g.shape[0] < templates[i].shape[0] or img_g.shape[1] < templates[i].shape[1]:
44 #         break
45
46 # result = cv2.matchTemplate(img_canny, templates[i], cv2.TM_CCORR_NORMED) # img_g.shape - template_g.
47 #     ↳ shape + 1
48 # _, max_val, _, max_loc = cv2.minMaxLoc(result)
49 # print(max_val)
50 # temp_img = img
51 # r = ratios[i]
52 # start_x, start_y = max_loc[0], max_loc[1]
53 # end_x, end_y = int((max_loc[0] + w * r)), int((max_loc[1] + h * r))
54 # cv2.rectangle(temp_img, (start_x, start_y), (end_x, end_y), (0, 255, 0), thickness=2)
55 # cv2.imshow("temp_img", temp_img)
56 # cv2.waitKey(0)
57
58 # if found is None or max_val > found[0]:
59 #     found = (max_val, max_loc, ratios[i])
60
61 # max_val, max_loc, r = found
62 # start_x, start_y = max_loc[0], max_loc[1]
63 # end_x, end_y = int((max_loc[0] + w * r)), int((max_loc[1] + h * r))
64
65 # cv2.rectangle(img, (start_x, start_y), (end_x, end_y), (0, 255, 0), thickness=2)
66 # cv2.imwrite(os.path.join(output_dir, img_name), img)
67
68 # cv2.imshow("img", img)
69 # cv2.waitKey(0)
70
71 def get_score(img1, img2):
72     sift = cv2.xfeatures2d.SIFT_create()
73     kp1, des1 = sift.detectAndCompute(img1, None)
74     kp2, des2 = sift.detectAndCompute(img2, None)
75     if len(kp1) < 2 or len(kp2) < 2:
76         return 0
77     index_params = dict(algorithm=0, trees=5)
78     flann = cv2.FlannBasedMatcher(index_params, None)
79     matches = flann.knnMatch(des1, des2, k=2)
80     good_matches = []
81     for m, n in matches:

```

```

81         if m.distance < 0.6*n.distance:
82             good_matches.append(m)
83     num_kps = 0
84     if len(kp1) <= len(kp2):
85         num_kps = len(kp1)
86     else:
87         num_kps = len(kp2)
88     score = len(good_matches) / num_kps * 100
89     return score
90
91 def logo_detection(input_dir, output_dir, logo_path, min_threshold):
92     print("detect_logo")
93
94     if not os.path.isdir(output_dir):
95         os.mkdir(output_dir)
96
97     template = cv2.imread(str(logo_path))
98     template_g = cv2.cvtColor(template, cv2.COLOR_BGR2GRAY)
99
100    w, h = template_g.shape[::-1]
101
102    # Compute template of different sizes
103    if template.shape[0] < 50:
104        scales = np.linspace(0.8, 1.0, 6)[::-1]
105    else:
106        # scales = np.linspace(0.1, 1.0, 25)[::-1]
107        scales = np.linspace(0.5, 1.0, 10)[::-1]
108    templates = []
109    ratios = []
110    for scale in scales:
111        resized = imutils.resize(template_g, width=int(template_g.shape[1] * scale))
112        template_canny = cv2.Canny(resized, 50, 200)
113        # templates.append(template_canny)
114        templates.append(resized)
115        ratios.append(resized.shape[1] / float(template_g.shape[1]))
116
117    for img_name in os.listdir(input_dir):
118        if not img_name.endswith(".jpg"):
119            continue
120        img = cv2.imread(os.path.join(input_dir, img_name))
121        img_g = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
122        img_canny = cv2.Canny(img_g, 100, 200)
123
124        p = -1
125        q = -1
126        matches = []
127        # Loop through the templates from small to big
128        for i in range(len(templates) - 1, -1, -1):
129            # Stop when template is bigger than image
130            if img_g.shape[0] < templates[i].shape[0] or img_g.shape[1] < templates[i].
131                ↳ shape[1]:
132                break
133
134            # match = cv2.matchTemplate(img_canny, templates[i], cv2.TM_CCORR_NORMED) # img_g.shape -
135                ↳ template_g.shape + 1
136            match = cv2.matchTemplate(img_g, templates[i], cv2.TM_CCORR_NORMED) # img_g.shape
137                ↳ - template_g.shape + 1
138            if p == -1 and q == -1:
139                p, q = match.shape

```

```

137     m, n = match.shape
138     matches.append(np.pad(match, ((0, p - m), (0, q - n)), mode="constant",
        ↪ constant_values=0))
139
140     boxes = []
141
142     matches = np.array(matches)
143     r, max_y, max_x = np.unravel_index(np.argmax(matches), matches.shape)
144     r = ratios[len(ratios) - 1 - r]
145     max_val = np.max(matches)
146     max_thresh = max(max_val * 0.95, min_threshold)
147     # print("max_val:", max_val, min_threshold, max_val < min_threshold)
148     if max_val < min_threshold:
149         cv2.imwrite(os.path.join(output_dir, img_name), img)
150         continue
151     start_x, start_y = max_x, max_y
152     end_x, end_y = int((max_x + w * r)), int((max_y + h * r))
153     match_score = get_score(template, img[start_y:end_y, start_x:end_x])
154     # if match_score > 0:
155         # print("How good it's the match: ", match_score, img_name)
156     if match_score < 5:
157         cv2.imwrite(os.path.join(output_dir, img_name), img)
158         continue
159     boxes.append((r, max_y, max_x, 1))
160
161     match_locations = np.where(matches >= max_thresh)
162     for i in range(len(match_locations[0])):
163         r1, y1, x1 = ratios[len(ratios) - 1 - match_locations[0][i]], match_locations
        ↪ [1][i], match_locations[2][i]
164         found = False
165         for j in range(len(boxes)):
166             r2, y2, x2, count = boxes[j]
167             # Check if two boxes of the same size overlap or
168             # if a smaller one is contained in the bigger one
169             if (r1 == r2 and np.abs(x1 - x2) < w * r1 and np.abs(y1 - y2) < h * r1) or
        ↪ \
170                 ((r1 < r2 and (x1 <= (x2 + w * r2) <= (x1 + w * r1)) and (y1 <= (y2 + h
        ↪ * r2) <= (y1 + h * r1))) or\
171                 (np.abs(x1 - x2) < 0.5 * w * r1 and np.abs((x1 + w * r1) - (x2 + w * r2
        ↪ )) < 0.5 * w * r1 and\
172                 np.abs(y1 - y2) < 0.5 * h * r1 and np.abs((y1 + h * r1) - (y2 + h * r2
        ↪ ) < 0.5 * h * r1):
173                 boxes[j] = (r2, y2, x2, count + 1)
174                 found = True
175                 break
176         if not found:
177             start_x, start_y = x1, y1
178             end_x, end_y = int((x1 + w * r1)), int((y1 + h * r1))
179             # cv2.imshow("", img[start_y:end_y, start_x:end_x])
180             # cv2.waitKey(0)
181             match_score = get_score(template, img[start_y:end_y, start_x:end_x])
182             # if match_score > 0:
183                 # print("How good it's the match: ", match_score, img_name)
184             if match_score > 5:
185                 boxes.append((r1, y1, x1, 1))
186
187     # print("boxes:", boxes)
188
189     for r, y, x, count in boxes:

```

```

190     start_x, start_y = x, y
191     end_x, end_y = int((x + w * r)), int((y + h * r))
192     # print(start_x, start_y, end_x, end_y)
193     cv2.rectangle(img, (start_x, start_y), (end_x, end_y), (0, 255, 0), thickness
        ↪ =2)
194
195     cv2.imwrite(os.path.join(output_dir, img_name), img)
196     # cv2.imshow("img", img)
197     # cv2.waitKey(0)

```

Listing 3: face.py

```

1  import copy
2  import cv2
3  import matplotlib.pyplot as plt
4  import numpy as np
5  import os
6  import pickle
7  import scipy.io, scipy.misc
8  from skimage import io
9  from skimage import transform as tf
10 from skimage.color import rgb2hsv
11 from sklearn import svm
12 from utils import shuffle
13
14 def visualize_distributions(imgs_dir):
15     """
16     Visualize the distributions of values in HSV and BGR of pictures in imgs_dir.
17     """
18     H, S, V = None, None, None
19     B, G, R = None, None, None
20     for filename in os.listdir(imgs_dir):
21         if not filename.endswith(".png"):
22             continue
23         img = cv2.imread(imgs_dir + filename)
24         hsv_img = cv2.cvtColor(img, cv2.COLOR_BGR2HSV)
25         if H is None:
26             H = hsv_img[..., 0].flatten()
27             S = hsv_img[..., 1].flatten()
28             V = hsv_img[..., 2].flatten()
29             B = img[..., 0].flatten()
30             G = img[..., 1].flatten()
31             R = img[..., 2].flatten()
32         else:
33             H = np.concatenate([H, hsv_img[..., 0].flatten()])
34             S = np.concatenate([S, hsv_img[..., 1].flatten()])
35             V = np.concatenate([V, hsv_img[..., 2].flatten()])
36             B = np.concatenate([B, img[..., 0].flatten()])
37             G = np.concatenate([G, img[..., 1].flatten()])
38             R = np.concatenate([R, img[..., 2].flatten()])
39
40     # Plot
41     f = plt.figure()
42     ax1 = f.add_subplot(1, 3, 1)
43     ax1.hist(H, bins=180,
44             range=(0.0, 180.0), histtype="stepfilled", color="b", label="Hue")
45     plt.title("Hue")
46     plt.xlabel("Value")
47     plt.ylabel("Frequency")
48     ax2 = f.add_subplot(1, 3, 2)

```

```

49     ax2.hist(S, bins=256,
50             range=(0.0, 255.0), histtype="stepfilled", color="g", label="Saturation")
51     plt.title("Saturation")
52     plt.xlabel("Value")
53     plt.ylabel("Frequency")
54     ax3 = f.add_subplot(1, 3, 3)
55     ax3.hist(V, bins=256,
56            range=(0.0, 255.0), histtype="stepfilled", color="r", label="Value")
57     plt.title("Value")
58     plt.xlabel("Value")
59     plt.ylabel("Frequency")
60     f.tight_layout()
61     f.savefig(imgs_dir + "_hsv_distributions.png")
62     plt.show()
63
64     f = plt.figure()
65     ax1 = f.add_subplot(1, 3, 1)
66     ax1.hist(B, bins=256,
67            range=(0.0, 255.0), histtype="stepfilled", color="b", label="Blue")
68     plt.title("Blue")
69     plt.xlabel("Value")
70     plt.ylabel("Frequency")
71     ax2 = f.add_subplot(1, 3, 2)
72     ax2.hist(G, bins=256,
73            range=(0.0, 255.0), histtype="stepfilled", color="g", label="Saturation")
74     plt.title("Green")
75     plt.xlabel("Value")
76     plt.ylabel("Frequency")
77     ax3 = f.add_subplot(1, 3, 3)
78     ax3.hist(R, bins=256,
79            range=(0.0, 255.0), histtype="stepfilled", color="r", label="Red")
80     plt.title("Red")
81     plt.xlabel("Value")
82     plt.ylabel("Frequency")
83     f.tight_layout()
84     f.savefig(imgs_dir + "_rgb_distributions.png")
85     plt.show()
86
87 def crop_images(old_dir, new_dir):
88     """
89     Crop .jpg images in old_dir given coordinates of
90     left eye, right eye, nose and mouth in .mat files.
91     Save cropped images in new_dir.
92     """
93     print("crop_images")
94     os.mkdir(new_dir)
95
96     H, S, V = None, None, None
97     for filename in os.listdir(old_dir):
98         if not filename.endswith(".jpg"):
99             continue
100
101         index = filename.find(".jpg")
102         name = filename[:index]
103
104         # Approximate coordinates of face
105         coords = scipy.io.loadmat(old_dir + name + ".mat")
106         start_x = int(coords["x"][0][0] - 0.5*(coords["x"][1][0] - coords["x"][0][0]))
107         end_x = int(coords["x"][1][0] + 0.5*(coords["x"][1][0] - coords["x"][0][0]))

```

```

108     start_y = int(coords["y"][0][0] - (coords["y"][3][0] - coords["y"][0][0]))
109     end_y = int(coords["y"][3][0] + (coords["y"][3][0] - coords["y"][2][0]))
110     img = io.imread(old_dir + filename)
111     face = img[start_y:end_y, start_x:end_x]
112     # Save cropped image
113     scipy.misc.imsave(new_dir + name + ".png", face)
114
115 def resize_images(input_dir, output_dir, size=72):
116     """
117     Resize images in input_dir to size x size.
118     """
119     print("resize_images")
120     if not os.path.isdir(output_dir):
121         os.mkdir(output_dir)
122
123     for filename in os.listdir(input_dir):
124         if not filename.endswith(".png"):
125             continue
126         img = cv2.imread(os.path.join(input_dir, filename))
127         x, y, d = img.shape
128         # Pad image to square
129         if x > y:
130             padded_img = np.pad(img, ((0, 0), (0, x - y), (0, 0)),
131                                 mode="constant", constant_values=0)
132         else:
133             padded_img = np.pad(img, ((0, y - x), (0, 0), (0, 0)),
134                                 mode="constant", constant_values=0)
135         resized_img = cv2.resize(padded_img, (size, size))
136         scipy.misc.imsave(output_dir + filename, cv2.cvtColor(resized_img, cv2.
137             ↪ COLOR_BGR2RGB))
138
139 def get_features(input_dir):
140     """
141     Get keypoints and descriptors of the images in input_dir with SIFT.
142     """
143     kps, des = None, None
144     sift = cv2.xfeatures2d.SIFT_create()
145
146     for filename in os.listdir(input_dir):
147         if not filename.endswith(".png"):
148             continue
149         img = io.imread(input_dir + filename)
150         kp, d = sift.detectAndCompute(img, None)
151
152         if des is None:
153             kps = kp
154             des = d
155         else:
156             kps = np.concatenate([kps, kp], axis=0)
157             des = np.concatenate([des, d], axis=0)
158     return kps, des
159
160 def get_data(f_dir, m_dir):
161     """
162     Get the images of the F and M classes from their respective directories and shuffle.
163     """
164     import keras
165     x = []
166     y = [0] * len(os.listdir(f_dir)) + [1] * len(os.listdir(m_dir))

```

```

166     imgs = os.listdir(f_dir) + os.listdir(m_dir)
167     for filename in os.listdir(f_dir):
168         if not filename.endswith(".png"):
169             continue
170         x.append(cv2.imread(os.path.join(f_dir, filename)))
171     for filename in os.listdir(m_dir):
172         if not filename.endswith(".png"):
173             continue
174         x.append(cv2.imread(os.path.join(m_dir, filename)))
175
176     x, y = shuffle(x, y)
177     y = keras.utils.to_categorical(y, num_classes=2)
178     return x, y
179
180 def train_model(model_path, classification):
181     OLD_F_DIR = "original_data/female/"
182     OLD_M_DIR = "original_data/male/"
183     F_TRAIN_DIR = "data/female_train/"
184     M_TRAIN_DIR = "data/male_train/"
185     F_TEST_DIR = "data/female_test/"
186     M_TEST_DIR = "data/male_test/"
187     F_TRAIN_CNN_DIR = "data/female_cnn_train/"
188     M_TRAIN_CNN_DIR = "data/male_cnn_train/"
189     F_TEST_CNN_DIR = "data/female_cnn_test/"
190     M_TEST_CNN_DIR = "data/male_cnn_test/"
191
192     if not os.path.isdir("data/"):
193         os.mkdir("data/")
194         crop_images(OLD_F_DIR, F_TRAIN_DIR)
195         crop_images(OLD_M_DIR, M_TRAIN_DIR)
196
197     if classification == "SVM":
198         f_train_kps, f_train_des = get_features(F_TRAIN_DIR)
199         m_train_kps, m_train_des = get_features(M_TRAIN_DIR)
200         x_train = np.concatenate([f_train_des, m_train_des], axis=0)
201         y_train = [-1] * len(f_train_des) + [1] * len(m_train_des)
202         model = svm.SVC(kernel="rbf", gamma="scale", C=10.0)
203         print("Start training")
204         model.fit(x_train, y_train)
205         # Save model
206         pickle.dump(model, open(model_path, "wb"))
207         print("Model saved as " + str(model_path))
208
209     elif classification == "NN_SIFT":
210         import keras
211         from keras.layers import Conv2D, Dense, Dropout, Flatten, MaxPooling2D
212         from keras.models import Sequential
213         from keras.optimizers import Adam, SGD
214         from keras.utils import plot_model
215
216         f_train_kps, f_train_des = get_features(F_TRAIN_DIR)
217         m_train_kps, m_train_des = get_features(M_TRAIN_DIR)
218         x_train = np.concatenate([f_train_des, m_train_des], axis=0)
219         y_train = [0] * len(f_train_des) + [1] * len(m_train_des)
220         x_train, y_train = shuffle(x_train, y_train)
221
222         f_test_kps, f_test_des = get_features(F_TEST_DIR)
223         m_test_kps, m_test_des = get_features(M_TEST_DIR)
224         x_test = np.concatenate([f_test_des, m_test_des], axis=0)

```

```

225     y_test = [0] * len(f_test_des) + [1] * len(m_test_des)
226     x_test, y_test = shuffle(x_test, y_test)
227
228     model = Sequential()
229
230     model.add(Dense(64, activation='relu', kernel_initializer='random_normal', input_dim
        ↪ =128))
231     model.add(Dense(64, activation='relu', kernel_initializer='random_normal'))
232     model.add(Dense(64, activation='relu', kernel_initializer='random_normal'))
233     model.add(Dense(64, activation='relu', kernel_initializer='random_normal'))
234     model.add(Dense(1, activation='sigmoid', kernel_initializer='random_normal'))
235
236     model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
237
238     model.fit(x_train, y_train, batch_size=10, epochs=20)
239     score = model.evaluate(x_test, y_test)
240     print("score:", score)
241
242     model.save(model_path)
243     print("Model saved as " + str(model_path))
244     plot_model(model, to_file="output/nn_model.png", show_shapes=True)
245
246 elif classification == "CNN":
247     import keras
248     from keras.layers import Conv2D, Dense, Dropout, Flatten, MaxPooling2D
249     from keras.models import Sequential
250     from keras.optimizers import Adam, SGD
251     from keras.utils import plot_model
252
253     np.random.seed(123)
254     size = 72
255
256     # resize_images(F_TRAIN_DIR, F_TRAIN_CNN_DIR, size)
257     # resize_images(M_TRAIN_DIR, M_TRAIN_CNN_DIR, size)
258     # resize_images(F_TEST_DIR, F_TEST_CNN_DIR, size)
259     # resize_images(M_TEST_DIR, M_TEST_CNN_DIR, size)
260
261     x_train, y_train = get_data(F_TRAIN_CNN_DIR, M_TRAIN_CNN_DIR)
262     x_test, y_test = get_data(F_TEST_CNN_DIR, M_TEST_CNN_DIR)
263
264     model = Sequential()
265     model.add(Conv2D(32, (3, 3), activation="relu", input_shape=(size, size, 3)))
266     model.add(Conv2D(32, (3, 3), activation="relu"))
267     model.add(MaxPooling2D(pool_size=(2, 2)))
268     model.add(Dropout(0.25))
269
270     model.add(Conv2D(64, (3, 3), activation="relu"))
271     model.add(Conv2D(64, (3, 3), activation="relu"))
272     model.add(MaxPooling2D(pool_size=(2, 2)))
273     model.add(Dropout(0.25))
274
275     model.add(Conv2D(128, (3, 3), activation="relu"))
276     model.add(Conv2D(128, (3, 3), activation="relu"))
277     model.add(MaxPooling2D(pool_size=(2, 2)))
278     model.add(Dropout(0.25))
279
280     model.add(Flatten())
281     model.add(Dense(256, activation="relu"))
282     model.add(Dropout(0.5))

```



```

283     model.add(Dense(2, activation="softmax"))
284
285     adam = Adam(lr=0.001)
286     model.compile(loss="binary_crossentropy", optimizer=adam, metrics=["accuracy"])
287
288     model.fit(x_train, y_train, epochs=10)
289     score = model.evaluate(x_test, y_test)
290     print("score:", score)
291
292     model.save(model_path)
293     print("Model saved as " + str(model_path))
294     plot_model(model, to_file="output/cnn_model.png", show_shapes=True)
295
296 else:
297     raise ValueError("Illegal classification value")
298
299 def predict_model(model_path):
300     print("predict_model")
301     F_TEST_DIR = "data/female_test/"
302     M_TEST_DIR = "data/male_test/"
303
304     ext = os.path.splitext(model_path)[1]
305     if ext == ".sav":
306         model = pickle.load(open(model_path, "rb"))
307     elif ext == ".h5":
308         from keras.models import load_model
309         model = load_model(model_path)
310     else:
311         raise ValueError("Not valid model extension")
312
313     sift = cv2.xfeatures2d.SIFT_create()
314
315     correct = 0
316     for filename in os.listdir(F_TEST_DIR):
317         if not filename.endswith(".png"):
318             continue
319         img = cv2.imread(os.path.join(F_TEST_DIR, filename))
320         kps, des = sift.detectAndCompute(img, None)
321         result = model.predict(des)
322         if ext == ".sav":
323             f_count = np.count_nonzero(result == -1)
324             m_count = np.count_nonzero(result == 1)
325         elif ext == ".h5":
326             f_count = np.sum(result < 0.5)
327             m_count = np.sum(result > 0.5)
328         if f_count > m_count:
329             correct += 1
330
331     for filename in os.listdir(M_TEST_DIR):
332         if not filename.endswith(".png"):
333             continue
334         img = cv2.imread(os.path.join(M_TEST_DIR, filename))
335         kps, des = sift.detectAndCompute(img, None)
336         result = model.predict(des)
337         if ext == ".sav":
338             f_count = np.count_nonzero(result == -1)
339             m_count = np.count_nonzero(result == 1)
340         elif ext == ".h5":
341             f_count = np.sum(result < 0.5)

```

```

342         m_count = np.sum(result > 0.5)
343         if m_count > f_count:
344             correct += 1
345         score = correct / (len(os.listdir(F_TEST_DIR)) + len(os.listdir(M_TEST_DIR))) * 100
346         print(str(score) + "% of images are categorized correctly")
347
348     def face_detection_hsv(input_dir, output_dir):
349         print("face_detection_hsv")
350
351         if not os.path.isdir(output_dir):
352             os.mkdir(output_dir)
353         for filename in os.listdir(input_dir):
354             if not filename.endswith(".jpg"):
355                 continue
356             img = cv2.imread(os.path.join(input_dir, filename))
357
358             # HSV color detection
359             hsv_img = cv2.cvtColor(img, cv2.COLOR_BGR2HSV)
360
361             lower1 = np.array([0, 48, 80])
362             upper1 = np.array([20, 255, 255])
363             mask1 = cv2.inRange(hsv_img, lower1, upper1) # img_hsv.shape
364             lower2 = np.array([170, 0, 0])
365             upper2 = np.array([180, 255, 255])
366             mask2 = cv2.inRange(hsv_img, lower2, upper2) # img_hsv.shape
367             mask = cv2.bitwise_or(mask1, mask2)
368
369             kernel = cv2.getStructuringElement(cv2.MORPH_ELLIPSE, (10, 10))
370             new_img = cv2.bitwise_and(img, img, mask=mask)
371             cv2.imwrite(os.path.join(output_dir, filename), new_img)
372
373     def face_detection_cascade(input_dir, output_dir, model_path, classification):
374         print("face_detection_cascade")
375         XML_FILENAME = "haarcascade_frontalface_default.xml"
376         F_TEXT = "Female"
377         M_TEXT = "Male"
378         O_TEXT = "Not sure"
379         F_COLOR = (0, 0, 255)
380         M_COLOR = (255, 0, 0)
381         O_COLOR = (0, 255, 0)
382
383         if classification == "SVM":
384             model = pickle.load(open(model_path, "rb"))
385         elif classification == "NN_SIFT":
386             from keras.models import load_model
387             model = load_model(model_path)
388         elif classification == "CNN":
389             from keras.models import load_model
390             model = load_model(model_path)
391             size = 72
392         else:
393             raise ValueError("Illegal classification value")
394
395         if not os.path.isdir(output_dir):
396             os.mkdir(output_dir)
397
398         exts = [".jpg", ".png"]
399         img_names = [img for img in os.listdir(input_dir) if img.endswith(tuple(exts))]
400         # Sort images by name in ascending order

```

```

401 img_names.sort(key=lambda img: int(''.join(filter(str.isdigit, img))))
402 sift = cv2.xfeatures2d.SIFT_create()
403 index_params = dict(algorithm=0, trees=5)
404 flann = cv2.FlannBasedMatcher(index_params, None)
405 face_count = 0
406 prev_faces = {}
407
408 for filename in img_names:
409     img = cv2.imread(os.path.join(input_dir, filename))
410     img_g = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
411
412     #####
413     ### Cascade face detection ###
414     #####
415     face_cascade = cv2.CascadeClassifier(XML_FILENAME)
416     faces = face_cascade.detectMultiScale(img_g, scaleFactor=1.3,
417     minNeighbors=3, minSize=(70, 70))
418     curr_faces = {}
419
420     for x, y, w, h in faces:
421         #####
422         ### Face tracking ###
423         #####
424         face = img[y:y+h, x:x+w]
425         kp1, des1 = sift.detectAndCompute(face, None)
426         found_face = False
427         for (x2, y2, w2, h2), index in prev_faces.items():
428             face2 = img[y2:y2+h, x2:x2+w]
429             kp2, des2 = sift.detectAndCompute(face2, None)
430             if len(kp1) < 2 or len(kp2) < 2:
431                 continue
432             matches = flann.knnMatch(des1, des2, k=2)
433             good_matches = []
434             for m, n in matches:
435                 if m.distance < 0.6*n.distance:
436                     good_matches.append(m)
437             num_kps = 0
438             if len(kp1) <= len(kp2):
439                 num_kps = len(kp1)
440             else:
441                 num_kps = len(kp2)
442             score = len(good_matches) / num_kps * 100
443             if score > 45:
444                 found_face = True
445                 face_number = index
446                 curr_faces[(x, y, w, h)] = index
447         if not found_face:
448             face_number = face_count
449             curr_faces[(x, y, w, h)] = face_count
450             face_count += 1
451
452     #####
453     ### Gender classification ###
454     #####
455     if classification == "SVM":
456         # Get SIFT descriptors
457         kps, des = sift.detectAndCompute(face, None)
458         # Predict female or male
459         result = model.predict(des)

```

```

460     f_count = np.count_nonzero(result == -1)
461     m_count = np.count_nonzero(result == 1)
462     # Plot color box
463     if f_count > m_count:
464         text = str(face_number) + " " + F_TEXT + ": " + \
465             format(f_count/(f_count + m_count)*100, ".2f") + "%"
466         color = F_COLOR
467     elif m_count > f_count:
468         text = str(face_number) + " " + M_TEXT + ": " + \
469             format(m_count/(f_count + m_count)*100, ".2f") + "%"
470         color = M_COLOR
471     else:
472         text = str(face_number) + " " + O_TEXT
473         color = O_COLOR
474
475     elif classification == "NN_SIFT":
476         assert w == h
477         kps, des = sift.detectAndCompute(face, None)
478         prediction = model.predict(des)
479         f_count = np.sum(prediction < 0.5)
480         m_count = np.sum(prediction > 0.5)
481         if f_count > m_count:
482             text = str(face_number) + " " + F_TEXT + ": " + \
483                 format(f_count/(f_count + m_count)*100, ".2f") + "%"
484             color = F_COLOR
485         elif m_count > f_count:
486             text = str(face_number) + " " + M_TEXT + ": " + \
487                 format(m_count/(f_count + m_count)*100, ".2f") + "%"
488             color = M_COLOR
489         else:
490             text = str(face_number) + " " + O_TEXT
491             color = O_COLOR
492
493     elif classification == "CNN":
494         assert w == h
495         face = cv2.resize(face, (size, size))
496         prediction = model.predict(np.array([face]), verbose=0)
497         if prediction == -1:
498             text = str(face_number) + " " + F_TEXT + ": " + \
499                 format(prediction[0]*100, ".2f") + "%"
500             color = F_COLOR
501         elif prediction == 1:
502             text = str(face_number) + " " + M_TEXT + ": " + \
503                 format(prediction[0]*100, ".2f") + "%"
504             color = M_COLOR
505         else:
506             text = str(face_number) + " " + O_TEXT
507             color = O_COLOR
508         cv2.rectangle(img, (x, y), (x+w, y+h), color, thickness=2)
509         cv2.putText(img, text, (x, y-10), color=color,
510             fontFace=cv2.FONT_HERSHEY_PLAIN, fontScale=1)
511
512     prev_faces = curr_faces
513     cv2.imwrite(os.path.join(output_dir, filename), img)
514     cv2.imshow("img", img)
515     cv2.waitKey(0)

```

Listing 4: utils.py

```

1 import cv2

```

```

2 import os
3 import numpy as np
4
5 def make_video(imgs_dir, vid_name, fps):
6     """
7     Make vid_name.mp4 using images in imgs_dir.
8     """
9     OUTPUT_DIR = "output/"
10    if not os.path.isdir(OUTPUT_DIR):
11        os.mkdir(OUTPUT_DIR)
12    exts = [".jpg", ".png"]
13    imgs = [img for img in os.listdir(imgs_dir) if img.endswith(tuple(exts))]
14    # Sort images by name in ascending order
15    imgs.sort(key=lambda img: int("".join(filter(str.isdigit, img))))
16    frame = cv2.imread(os.path.join(imgs_dir, imgs[0]))
17    h, w, _ = frame.shape
18
19    # Make sure vid_name does not include a file extension
20    index = vid_name.find(".")
21    if index != -1:
22        vid_name = vid_name[:index]
23    file_path = OUTPUT_DIR + vid_name + ".mp4"
24    vid = cv2.VideoWriter(file_path, cv2.VideoWriter_fourcc(*"MP4V"), fps, (w, h))
25    for img in imgs:
26        vid.write(cv2.imread(os.path.join(imgs_dir, img)))
27    vid.release()
28    print("Video is now in ", file_path)
29
30 def shuffle(x, y):
31     """
32     Shuffle data x and their labels y.
33     """
34    assert len(x) == len(y)
35    idx = np.random.permutation(len(x))
36    x, y = np.array(x)[idx], np.array(y)[idx]
37    return x, y

```